

1 Start coding or [generate](#) with AI.

```
1 import os
2 import numpy as np
3 import pandas as pd
4 import cv2
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from PIL import Image
8
9 import tensorflow as tf
10 from tensorflow import keras
11 from tensorflow.keras import layers, models
12 from tensorflow.keras.models import Model
13 from tensorflow.keras.layers import Conv2D, MaxPooling2D, UpSampling2D, Input
14 from tensorflow.keras.layers import BatchNormalization, Concatenate
15 import tensorflow.keras.backend as K
16 from tensorflow.keras.preprocessing.image import ImageDataGenerator
17 from tensorflow.keras.layers import Conv2D, BatchNormalization, Activation, MaxPooling2D, UpSampling2D, Concatenate
18 from sklearn.metrics import mean_squared_error
```

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

1. Brief description of the problem and data

Description

In this project, we aim to develop a machine learning model to remove noise from scanned text documents and restore them to their clean, readable versions. Optical Character Recognition (OCR) is widely used to digitize printed and handwritten documents, but real-world issues like stains, wrinkles, and smudges make OCR less effective. The goal of this project is to create a document enhancement model that improves text clarity by removing noise while preserving the original content.

The methodology will include the following stages:

- Brief Description of the Problem and Data
- Data Preprocessing and Exploratory Data Analysis (EDA)
- Model Architecture and Development
- Model Evaluation
- Results and Analysis
- Conclusion

Data

The dataset consists of scanned text images with synthetic noise added to simulate real-world document degradation. It contains:

- train/: Noisy images used for training.
- train_cleaned/: Corresponding clean (ground truth) versions of the images.
- test/: Noisy test images where the model must remove noise.
- sampleSubmission.csv: A sample submission file that shows the required format for Kaggle.

Each image is grayscale, with pixel intensity values ranging from 0 (black) to 1 (white). The submission format requires flattening the image into a list of pixel values, each assigned a unique ID (image_row_col).

```
1 # Define dataset paths
2 dataset_path = "/content/drive/My Drive/D TSA5510/final/"
3 train_path = os.path.join(dataset_path, "train")
4 train_cleaned_path = os.path.join(dataset_path, "train_cleaned")
5 test_path = os.path.join(dataset_path, "test")

1 # Count images each folders
2 def count_images(directory):
3     return len(os.listdir(directory))
4
5 print("Number of Train Images:", count_images(train_path))
6 print("Number of Train Cleaned Images:", count_images(train_cleaned_path))
7 print("Number of Test Images:", count_images(test_path))

Number of Train Images: 144
Number of Train Cleaned Images: 144
Number of Test Images: 72

1 # Inspect
2 print("Train set sample files:", os.listdir(train_path)[:5])
3 print("Train Cleaned set sample files:", os.listdir(train_cleaned_path)[:5])
4 print("Test set sample files:", os.listdir(test_path)[:5])

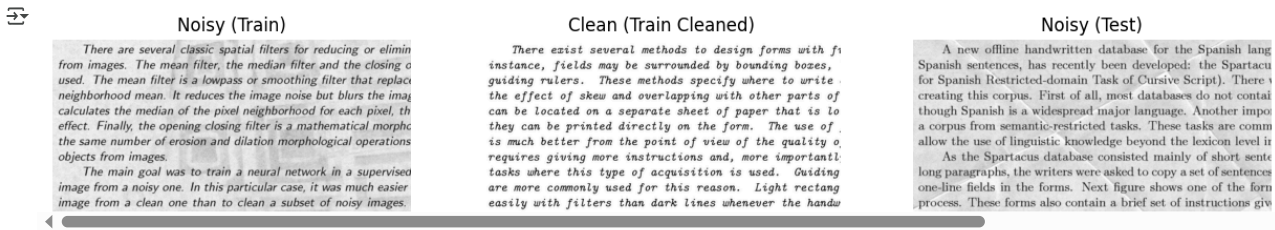
Train set sample files: ['33.png', '8.png', '42.png', '45.png', '15.png']
Train Cleaned set sample files: ['56.png', '36.png', '60.png', '47.png', '8.png']
Test set sample files: ['16.png', '40.png', '79.png', '91.png', '13.png']

1 # Inspect and display images
2 sample_noisy = os.path.join(train_path, os.listdir(train_path)[0])
3 sample_clean = os.path.join(train_cleaned_path, os.listdir(train_cleaned_path)[0])
4 sample_test = os.path.join(test_path, os.listdir(test_path)[0])
5
6 def display_images(img_paths, titles, figsize=(15, 5)):
7     fig, axes = plt.subplots(1, len(img_paths), figsize=figsize)
```

```

8     for ax, img_path, title in zip(axes, img_paths, titles):
9         img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
10        ax.imshow(img, cmap='gray')
11        ax.set_title(title)
12        ax.axis("off")
13    plt.show()
14
15 display_images(
16     [sample_noisy, sample_clean, sample_test],
17     ["Noisy (Train)", "Clean (Train Cleaned)", "Noisy (Test)"]
18 )

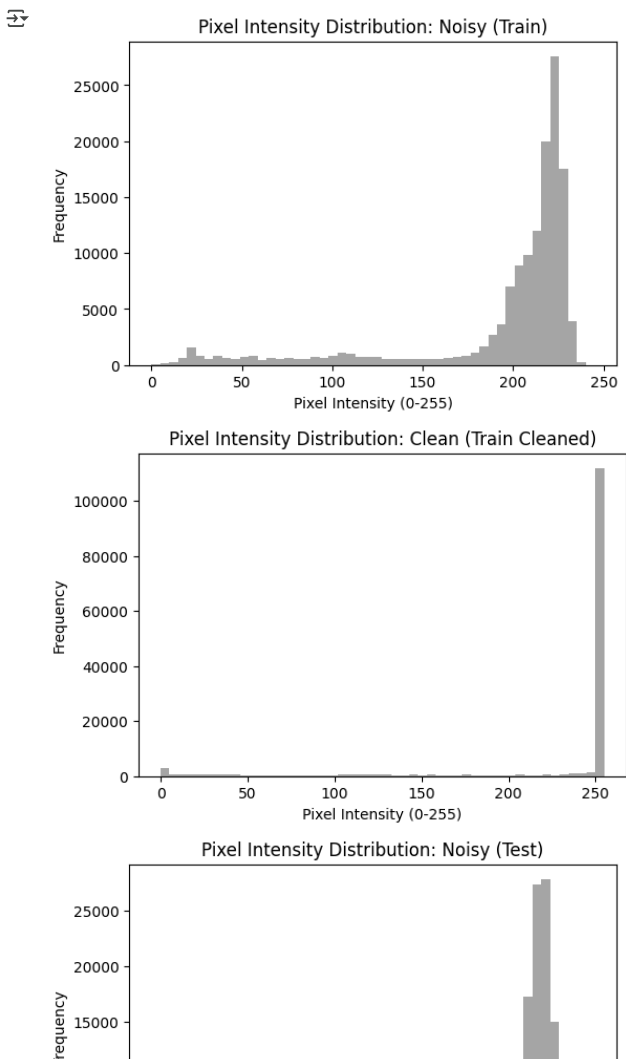
```



```

1 # Analyze Pixel Intensity Distribution
2 def plot_histogram(img_path, title):
3     img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
4     plt.figure(figsize=(6, 4))
5     plt.hist(img.ravel(), bins=50, color='gray', alpha=0.7)
6     plt.title(f"Pixel Intensity Distribution: {title}")
7     plt.xlabel("Pixel Intensity (0-255)")
8     plt.ylabel("Frequency")
9     plt.show()
10
11 # Plot histograms for noisy, cleaned, and test images
12 plot_histogram(sample_noisy, "Noisy (Train)")
13 plot_histogram(sample_clean, "Clean (Train Cleaned)")
14 plot_histogram(sample_test, "Noisy (Test)")

```



```

1 # Image Size Consistency
2 def check_image_shapes(directory):
3     shapes = []
4     for img_name in os.listdir(directory):
5         img_path = os.path.join(directory, img_name)

```

```

6         img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
7         shapes.append(img.shape)
8     return set(shapes)
9
10 # Check consistency in all datasets
11 print("Train Image Shapes:", check_image_shapes(train_path))
12 print("Train Cleaned Image Shapes:", check_image_shapes(train_cleaned_path))
13 print("Test Image Shapes:", check_image_shapes(test_path))

```

```

Train Image Shapes: {(420, 540), (258, 540)}
Train Cleaned Image Shapes: {(420, 540), (258, 540)}
Test Image Shapes: {(420, 540), (258, 540)}

```

EDA Summary

- **Dataset Completeness:** The dataset contains 144 training images with corresponding clean versions and 72 test images, ensuring full data availability for training and evaluation.
- **Image Size Variability:** The dataset has two unique image sizes, (420, 540) and (258, 540), requiring resizing for model consistency.
- **Noise Characteristics:** Noisy images contain dark artifacts, stains, and wrinkles, causing pixel intensity distributions to be more spread out compared to the clean images, which are mostly concentrated near 255 (white background).
- **Train vs. Test Similarity:** The test images exhibit similar noise patterns as the training set, confirming that denoising methods trained on the dataset should generalize well to test images.
- **Preprocessing Needs:** Due to varied noise types and intensity distributions, preprocessing should include grayscale conversion, normalization, and potential adaptive thresholding before applying denoising techniques.

3. Model Architecture and Development

Data Processing: Normalize pixel values and resize.

```

1 # Set image dimensions
2 IMG_HEIGHT, IMG_WIDTH = 256, 256

1 # Data Processing Functions
2 def load_images(directory, resize_shape=(IMG_HEIGHT, IMG_WIDTH)):
3     images = []
4     for img_name in sorted(os.listdir(directory)):
5         img_path = os.path.join(directory, img_name)
6         img = cv2.imread(img_path, cv2.IMREAD_GRAYSCALE)
7         img = cv2.resize(img, resize_shape)
8         img = img.astype("float32") / 255.0
9         images.append(img)
10    return np.array(images).reshape(-1, resize_shape[0], resize_shape[1], 1)

1 # Load noisy and clean images
2 train_noisy = load_images(train_path)
3 train_clean = load_images(train_cleaned_path)
4
5 # Print dataset shape
6 print(f"Train Noisy Shape: {train_noisy.shape}")
7 print(f"Train Clean Shape: {train_clean.shape}")

```

```

Train Noisy Shape: (144, 256, 256, 1)
Train Clean Shape: (144, 256, 256, 1)

```

Define Convolutional Autoencoder Model

```

1 # Define the autoencoder architecture
2 def build_autoencoder(input_shape=(IMG_HEIGHT, IMG_WIDTH, 1)):
3     input_img = Input(shape=input_shape)
4
5     # Encoder
6     x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
7     x = MaxPooling2D((2, 2), padding='same')(x)
8     x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
9     x = MaxPooling2D((2, 2), padding='same')(x)
10
11    # Decoder
12    x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
13    x = UpSampling2D((2, 2))(x)
14    x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
15    x = UpSampling2D((2, 2))(x)
16    output_img = Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x) # Output layer
17
18    # Compile model
19    autoencoder = Model(input_img, output_img)
20    autoencoder.compile(optimizer='adam', loss='mse')
21
22    return autoencoder

1 # Instantiate the model
2 autoencoder = build_autoencoder()

1 # Print model summary
2 autoencoder.summary()

```

Model: "functional"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 256, 256, 1)	0
conv2d (Conv2D)	(None, 256, 256, 32)	320
max_pooling2d (MaxPooling2D)	(None, 128, 128, 32)	0
conv2d_1 (Conv2D)	(None, 128, 128, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_2 (Conv2D)	(None, 64, 64, 64)	36,928
up_sampling2d (UpSampling2D)	(None, 128, 128, 64)	0
conv2d_3 (Conv2D)	(None, 128, 128, 32)	18,464
up_sampling2d_1 (UpSampling2D)	(None, 256, 256, 32)	0
conv2d_4 (Conv2D)	(None, 256, 256, 1)	289

Total params: 74,497 (291.00 KB)

Train the Model

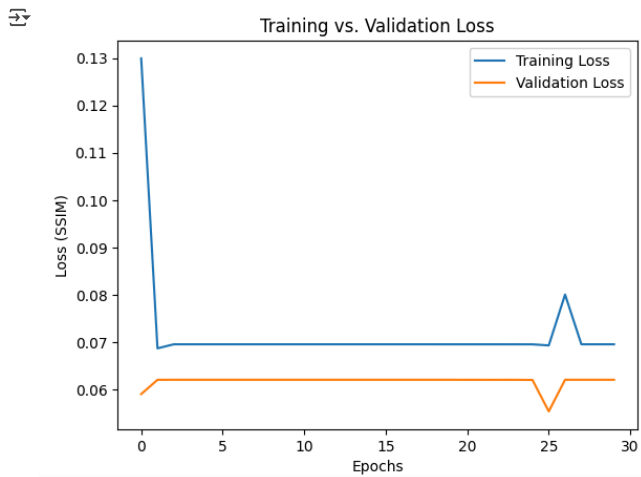
```
1 # Train the model
2 history = autoencoder.fit(
3     train_noisy, train_clean,
4     epochs=30, batch_size=16,
5     shuffle=True, validation_split=0.1
6 )
```

Epoch 1/30
9/9 11s 452ms/step - loss: 0.1628 - val_loss: 0.0590
Epoch 2/30
9/9 0s 42ms/step - loss: 0.0677 - val_loss: 0.0621
Epoch 3/30
9/9 0s 42ms/step - loss: 0.0697 - val_loss: 0.0621
Epoch 4/30
9/9 0s 42ms/step - loss: 0.0700 - val_loss: 0.0621
Epoch 5/30
9/9 0s 42ms/step - loss: 0.0696 - val_loss: 0.0621
Epoch 6/30
9/9 0s 42ms/step - loss: 0.0684 - val_loss: 0.0621
Epoch 7/30
9/9 0s 42ms/step - loss: 0.0694 - val_loss: 0.0621
Epoch 8/30
9/9 0s 42ms/step - loss: 0.0691 - val_loss: 0.0621
Epoch 9/30
9/9 0s 42ms/step - loss: 0.0694 - val_loss: 0.0621
Epoch 10/30
9/9 0s 42ms/step - loss: 0.0701 - val_loss: 0.0621
Epoch 11/30
9/9 0s 41ms/step - loss: 0.0698 - val_loss: 0.0621
Epoch 12/30
9/9 0s 42ms/step - loss: 0.0699 - val_loss: 0.0621
Epoch 13/30
9/9 0s 42ms/step - loss: 0.0704 - val_loss: 0.0621
Epoch 14/30
9/9 0s 42ms/step - loss: 0.0700 - val_loss: 0.0621
Epoch 15/30
9/9 0s 42ms/step - loss: 0.0712 - val_loss: 0.0621
Epoch 16/30
9/9 0s 42ms/step - loss: 0.0690 - val_loss: 0.0621
Epoch 17/30
9/9 0s 42ms/step - loss: 0.0702 - val_loss: 0.0621
Epoch 18/30
9/9 0s 42ms/step - loss: 0.0701 - val_loss: 0.0621
Epoch 19/30
9/9 0s 43ms/step - loss: 0.0693 - val_loss: 0.0621
Epoch 20/30
9/9 0s 42ms/step - loss: 0.0700 - val_loss: 0.0621
Epoch 21/30
9/9 0s 42ms/step - loss: 0.0714 - val_loss: 0.0621
Epoch 22/30
9/9 0s 44ms/step - loss: 0.0698 - val_loss: 0.0621
Epoch 23/30
9/9 0s 42ms/step - loss: 0.0696 - val_loss: 0.0621
Epoch 24/30
9/9 0s 42ms/step - loss: 0.0700 - val_loss: 0.0621
Epoch 25/30
9/9 0s 42ms/step - loss: 0.0703 - val_loss: 0.0620
Epoch 26/30
9/9 0s 42ms/step - loss: 0.0685 - val_loss: 0.0554
Epoch 27/30
9/9 0s 42ms/step - loss: 0.0867 - val_loss: 0.0621
Epoch 28/30
9/9 0s 42ms/step - loss: 0.0698 - val_loss: 0.0621
Epoch 29/30
9/9 0s 42ms/step - loss: 0.0698 - val_loss: 0.0621

```

1 # Plot training & validation loss
2 plt.plot(history.history['loss'], label='Training Loss')
3 plt.plot(history.history['val_loss'], label='Validation Loss')
4 plt.xlabel("Epochs")
5 plt.ylabel("Loss (SSIM)")
6 plt.legend()
7 plt.title("Training vs. Validation Loss")
8 plt.show()

```



Evaluate Model Performance

```

1 # Predict
2 test_noisy = load_images(test_path)
3 denoised_images = autoencoder.predict(test_noisy)

```

3/3 ————— 4s 494ms/step

```

1 # Display results
2 def display_denoising_results(noisy, denoised, num_samples=3):
3     fig, axes = plt.subplots(num_samples, 2, figsize=(10, 10))
4
5     for i in range(num_samples):
6         axes[i, 0].imshow(noisy[i].reshape(IMG_HEIGHT, IMG_WIDTH), cmap='gray')
7         axes[i, 0].set_title("Noisy Image")
8         axes[i, 0].axis("off")
9
10        axes[i, 1].imshow(denoised[i].reshape(IMG_HEIGHT, IMG_WIDTH), cmap='gray')
11        axes[i, 1].set_title("Denoised Image")
12        axes[i, 1].axis("off")
13
14    plt.show()

```

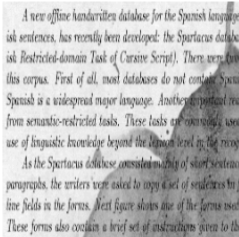
```

1 # Show test results
2 display_denoising_results(test_noisy, denoised_images)

```



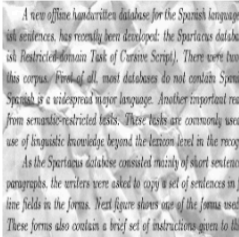
Noisy Image



Denoised Image




Noisy Image



Denoised Image



Noisy Image



Denoised Image



```

1 # Because we don't have test ground truth - So, I calculate RMSE based on Train data set to measure the performance.
2 denoised_train = autoencoder.predict(train_noisy)
3 rmse_scores = [
4     np.sqrt(mean_squared_error(train_clean[i].flatten(), denoised_train[i].flatten()))
5     for i in range(len(train_clean))
6 ]
7 overall_rmse = np.mean(rmse_scores)
8 overall_rmse

```



5/5 1s 264ms/step
Overall RMSE Score: 0.26120

The first model, the performance is good. The RMSE is 0.26 which is okay. However, we still can improve the denoised images which appear blurred and less clear, which suggests that the model is:

- Removing too much detail, affecting text sharpness.
- Smoothing out noise, but at the cost of reducing readability.
- Not preserving fine edges, leading to loss of text structure.

SSIM Loss Function: SSIM (Structural Similarity Index) measures how similar two images are, focusing on text clarity rather than pixel-wise differences.

```

1 def ssim_loss(y_true, y_pred):
2     return 1 - tf.reduce_mean(tf.image.ssim(y_true, y_pred, max_val=1.0))

1 # Define the autoencoder architecture (same as before)
2 def build_autoencoder(input_shape=(IMG_HEIGHT, IMG_WIDTH, 1)):
3     input_img = Input(shape=input_shape)
4
5     # Encoder
6     x = Conv2D(32, (3, 3), activation='relu', padding='same')(input_img)
7     x = MaxPooling2D((2, 2), padding='same')(x)
8     x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
9     x = MaxPooling2D((2, 2), padding='same')(x)
10
11    # Decoder
12    x = Conv2D(64, (3, 3), activation='relu', padding='same')(x)
13    x = UpSampling2D((2, 2))(x)
14    x = Conv2D(32, (3, 3), activation='relu', padding='same')(x)
15    x = UpSampling2D((2, 2))(x)
16    output_img = Conv2D(1, (3, 3), activation='sigmoid', padding='same')(x) # Output layer
17
18    # Compile model with SSIM loss
19    autoencoder = Model(input_img, output_img)
20    autoencoder.compile(optimizer='adam', loss=ssim_loss)
21
22    return autoencoder

```

```

1 # Build the autoencoder
2 enhance_autoencoder = build_autoencoder()
3 enhance_autoencoder.summary()

```

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 256, 256, 1)	0
conv2d_5 (Conv2D)	(None, 256, 256, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 128, 128, 32)	0
conv2d_6 (Conv2D)	(None, 128, 128, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 64, 64, 64)	0
conv2d_7 (Conv2D)	(None, 64, 64, 64)	36,928
up_sampling2d_2 (UpSampling2D)	(None, 128, 128, 64)	0
conv2d_8 (Conv2D)	(None, 128, 128, 32)	18,464
up_sampling2d_3 (UpSampling2D)	(None, 256, 256, 32)	0
conv2d_9 (Conv2D)	(None, 256, 256, 1)	289

Total params: 74,497 (291.00 KB)

```

1 # Train the model with SSIM loss
2 history = enhance_autoencoder.fit(
3     train_noisy, train_clean,
4     epochs=100,
5     batch_size=16,
6     shuffle=True,
7     validation_split=0.1
8 )

```

```

Epoch 1/100
9/9 ----- 6s 316ms/step - loss: 0.7918 - val_loss: 0.6993
Epoch 2/100
9/9 ----- 0s 46ms/step - loss: 0.7018 - val_loss: 0.5949
Epoch 3/100
9/9 ----- 0s 46ms/step - loss: 0.5820 - val_loss: 0.5473
Epoch 4/100
9/9 ----- 0s 46ms/step - loss: 0.5467 - val_loss: 0.5300
Epoch 5/100
9/9 ----- 0s 47ms/step - loss: 0.5259 - val_loss: 0.5177
Epoch 6/100
9/9 ----- 0s 46ms/step - loss: 0.5032 - val_loss: 0.4936
Epoch 7/100
9/9 ----- 0s 47ms/step - loss: 0.4812 - val_loss: 0.4688
Epoch 8/100
9/9 ----- 0s 46ms/step - loss: 0.4557 - val_loss: 0.4542
Epoch 9/100
9/9 ----- 0s 46ms/step - loss: 0.4413 - val_loss: 0.4501
Epoch 10/100
9/9 ----- 0s 46ms/step - loss: 0.4258 - val_loss: 0.4228
Epoch 11/100
9/9 ----- 0s 46ms/step - loss: 0.4091 - val_loss: 0.4213
Epoch 12/100
9/9 ----- 0s 46ms/step - loss: 0.3976 - val_loss: 0.4056
Epoch 13/100
9/9 ----- 0s 46ms/step - loss: 0.3847 - val_loss: 0.4082
Epoch 14/100
9/9 ----- 0s 47ms/step - loss: 0.3778 - val_loss: 0.3968
Epoch 15/100
9/9 ----- 0s 47ms/step - loss: 0.3721 - val_loss: 0.3770
Epoch 16/100
9/9 ----- 0s 47ms/step - loss: 0.3582 - val_loss: 0.3761
Epoch 17/100
9/9 ----- 0s 46ms/step - loss: 0.3579 - val_loss: 0.3658
Epoch 18/100
9/9 ----- 0s 46ms/step - loss: 0.3423 - val_loss: 0.3595
Epoch 19/100
9/9 ----- 0s 46ms/step - loss: 0.3372 - val_loss: 0.3528
Epoch 20/100
9/9 ----- 0s 46ms/step - loss: 0.3358 - val_loss: 0.3532
Epoch 21/100
9/9 ----- 0s 47ms/step - loss: 0.3277 - val_loss: 0.3482
Epoch 22/100
9/9 ----- 0s 47ms/step - loss: 0.3197 - val_loss: 0.3346
Epoch 23/100
9/9 ----- 0s 46ms/step - loss: 0.3161 - val_loss: 0.3354
Epoch 24/100
9/9 ----- 0s 47ms/step - loss: 0.3075 - val_loss: 0.3229
Epoch 25/100
9/9 ----- 0s 46ms/step - loss: 0.3047 - val_loss: 0.3133
Epoch 26/100
9/9 ----- 0s 46ms/step - loss: 0.2897 - val_loss: 0.3119
Epoch 27/100
9/9 ----- 0s 46ms/step - loss: 0.2831 - val_loss: 0.2965
Epoch 28/100
9/9 ----- 0s 46ms/step - loss: 0.2726 - val_loss: 0.2856
Epoch 29/100
9/9 ----- 0s 46ms/step - loss: 0.2571 - val_loss: 0.2775


```

```

1 # Plot training & validation loss
2 plt.plot(history.history['loss'], label='Training Loss')
3 plt.plot(history.history['val_loss'], label='Validation Loss')
4 plt.xlabel("Epochs")
5 plt.ylabel("Loss (SSIM)")
6 plt.legend()

```

```
7 plt.title("Training vs. Validation Loss")
8 plt.show()
```

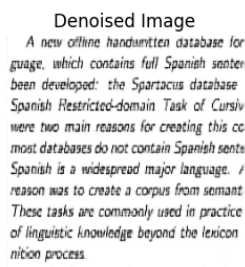


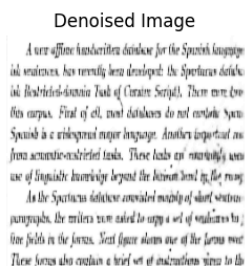
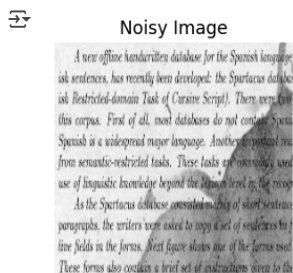


```
1 # Predict
2 test_noisy = load_images(test_path)
3 enhance_denoised_images = enhance_autoencoder.predict(test_noisy)
```

3/3 1s 169ms/step

```
1 # Display results
2 def display_denoising_results(noisy, denoised, num_samples=3):
3     fig, axes = plt.subplots(num_samples, 2, figsize=(10, 10))
4
5     for i in range(num_samples):
6         axes[i, 0].imshow(noisy[i].reshape(IMG_HEIGHT, IMG_WIDTH), cmap='gray')
7         axes[i, 0].set_title("Noisy Image")
8         axes[i, 0].axis("off")
9
10        axes[i, 1].imshow(denoised[i].reshape(IMG_HEIGHT, IMG_WIDTH), cmap='gray')
11        axes[i, 1].set_title("Denoised Image")
12        axes[i, 1].axis("off")
13
14    plt.show()
```


```
1 # Show test results
2 display_denoising_results(test_noisy, enhance_denoised_images)
```




```

1 # Because we don't have test ground truth - so, I calculate RMSE based on train
  data set to measure the performance.
2 denoised_train = enhance_autoencoder.predict(train_noisy)
3 rmse_scores = [
4     np.sqrt(mean_squared_error(train_clean[i].flatten(), denoised_train[i].
      flatten()))
5     for i in range(len(train_clean))
6 ]
7
8 overall_rmse = np.mean(rmse_scores)
9
10 print(f"Overall RMSE Score: {overall_rmse:.5f}")

```

5/5  0s 66ms/step
Overall RMSE Score: 0.11230

Results and Analysis

Please see the comparison between original, denoised with autoencoder and denoised with enhanced autoencoder model below.

- Convolutional Autoencoder performs well with SSIM loss and longer training, Enhanced Autoencoder. The denoised images show clear improvements in quality. It produces sharper, more readable text compared to the basic autoencoder, effectively reducing noise and improving contrast.
- The RMSE confirms that the enhancement model is better, RMSE 0.1123 comparing to RMSE, 0.26120 from previous model.
- The basic autoencoder struggles with fine text details, producing blurred outputs, while the enhanced model preserves more text structures and improves readability.
- Future improvements could explore GAN-based denoising or Transformer-based architectures to further enhance the model's ability to restore text clarity.

```

1 # Select random indices to visualize
2 num_images = 3
3 random_indices = np.random.choice(len(test_noisy), num_images, replace=False)
4
5 # Plot comparison
6 fig, axes = plt.subplots(num_images, 3, figsize=(10, 15))
7
8 for i, idx in enumerate(random_indices):
9     # Original test image
10    test_img = test_noisy[idx]
11    denoised_img = denoised_images[idx]
12    enhanced_img = enhance_denoised_images[idx]
13
14    # Plot original
15    axes[i, 0].imshow(test_img, cmap='gray')
16    axes[i, 0].set_title("Original Test Image")
17    axes[i, 0].axis("off")
18
19    # Plot autoencoder denoised image
20    axes[i, 1].imshow(denoised_img, cmap='gray')
21    axes[i, 1].set_title("Denoised (Autoencoder)")
22    axes[i, 1].axis("off")
23
24    # Plot enhanced autoencoder denoised image
25    axes[i, 2].imshow(enhanced_img, cmap='gray')
26    axes[i, 2].set_title("Denoised (Enhanced Autoencoder)")
27    axes[i, 2].axis("off")
28
29 plt.tight_layout()
30 plt.show()

```

Original Test Image

A new offline handwritten database for the which contains full Spanish sentences, has re the Spartacus database (which stands for Spar Task of Cursive Script). There were two main this corpus. First of all, most databases do sentences, even though Spanish is a widespread important reason was to create a corpus from tasks. These tasks are commonly used in practice of linguistic knowledge beyond the lexicon process.

As the Spartacus database consisted mainly and did not contain long paragraphs, the writer a set of sentences in fixed places: dedicated the forms. Next figure shows one of the forms process. These forms also contain a brief summary

Denoised (Autoencoder)

A new offline handwritten database for the which contains full Spanish sentences, has re the Spartacus database (which stands for Spar Task of Cursive Script). There were two main this corpus. First of all, most databases do sentences, even though Spanish is a widespread important reason was to create a corpus from tasks. These tasks are commonly used in practice of linguistic knowledge beyond the lexicon process.

As the Spartacus database consisted mainly and did not contain long paragraphs, the writer a set of sentences in fixed places: dedicated the forms. Next figure shows one of the forms process. These forms also contain a brief summary

Denoised (Enhanced Autoencoder)

A new offline handwritten database for the which contains full Spanish sentences, has re the Spartacus database (which stands for Spar Task of Cursive Script). There were two main this corpus. First of all, most databases do sentences, even though Spanish is a widespread important reason was to create a corpus from tasks. These tasks are commonly used in practice of linguistic knowledge beyond the lexicon process.

As the Spartacus database consisted mainly and did not contain long paragraphs, the writer a set of sentences in fixed places: dedicated the forms. Next figure shows one of the forms process. These forms also contain a brief summary